

# **Reinforcement Learning**

Lecture 4 - Model-free prediction

Per Mattsson

2022

Department of Information Technology

Repetition

### Important concepts



- States, actions and rewards:  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}$ ,  $r \in \mathcal{R}$ .
- Dynamics/model: p(s', r|s, a).
- **Policy:**  $\pi(a|s)$  (For deterministic policy also  $a = \pi(s)$ .)
- The return:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots$$

Is in general a stochastic variable.

• State-value function:

Expected return when starting in s and following policy  $\pi$ ,

$$v_{\pi}(s) = \mathbb{E}_{\pi}\left[G_t|S_t = s\right].$$

Action-value function:

Expected return when starting in s, taking action a and then follow  $\pi$ ,

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}\left[G_t|S_t = s, A_t = a\right]$$

# **Policy iteration**



Bellman equation:

$$v_{\pi}(s) = \sum_{\mathbf{a}} \pi(\mathbf{a}|s) \sum_{s',r} p(s',r|s,\mathbf{a})[r + \gamma v_{\pi}(s')], \quad \text{for all } s \in \mathcal{S}.$$

Policy evaluation:

$$v_{k+1}(s) = \sum_{\mathbf{a}} \pi(\mathbf{a}|s) \sum_{s',r} p(s',r|s,\mathbf{a})[r + \gamma v_k(s')]$$

then  $v_k(s) \rightarrow v_{\pi}(s)$ .

• Policy improvement:

$$q_{\pi}(s, \mathbf{a}) = \sum_{s', r} p(s', r|s, \mathbf{a})[r + \gamma v_{\pi}(s')]$$
$$\pi'(s) = \arg\max_{\mathbf{a}} q_{\pi}(s, \mathbf{a})$$

• Policy iteration:  $\pi_0 \xrightarrow{\mathsf{E}} \nu_{\pi_0} \xrightarrow{\mathsf{I}} \pi_1 \xrightarrow{\mathsf{E}} \nu_{\pi_1} \xrightarrow{\mathsf{I}} \pi_2 \xrightarrow{\mathsf{E}} \cdots \xrightarrow{\mathsf{I}} \pi_* \xrightarrow{\mathsf{E}} \nu_*.$ 

#### Value iteration



Bellman optimality equation:

$$v_*(s) = \max_{\mathbf{a}} q_*(s, \mathbf{a}) = \max_{\mathbf{a}} \sum_{r,s'} p(s', r|s, \mathbf{a})[r + \gamma v_*(s')].$$

• Value iteration: (Based on Bellman optimality equation)

$$v_{k+1}(s) = \max_{\mathbf{a}} \sum_{s',r} p(s',r|s,\mathbf{a})[r + \gamma v_k(s')]$$

then  $v_k(s) \rightarrow v_*(s)$ .

Optimal policy:

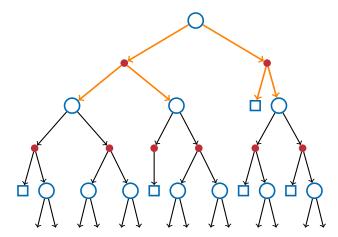
$$q_*(s, \mathbf{a}) = \sum_{r,s'} p(s', r|s, \mathbf{a})[r + \gamma v_*(s')]$$
 $\pi_*(s) = \arg\max_{\mathbf{a}} q_*(s, \mathbf{a}).$ 

# **Dynamic Programming Backup Diagram**



$$V(s) \leftarrow \mathbb{E}_{\pi} \left[ R_{t+1} + \gamma V(S_{t+1}) | S_t = s \right]$$

We need p(s', r|s, a) to compute the expected value.



# Monte-Carlo Methods

# **Example: Expected sum of throw with two dice**



- **Problem:** We throw two dice, and call their sum G. What is  $V = \mathbb{E}[G]$ ?
- By hand:
  - Each dice has 6 sides, so we can get  $6 \times 6 = 36$  combinations.
  - There is no way to get G = 1, so p(1) = 0.
  - There is one way to get G = 2, so p(2) = 1/36.
  - There are two ways to get G = 3, so p(3) = 2/36.
  - Etc.
  - Finally,  $\mathbb{E}[G] = \sum_{g=1}^{12} gp(g) = 7$ .
- Monte-Carlo:
  - Throw many times and get independent observations  $G_1, G_2, G_3, \ldots, G_n$ .
  - Use the empirical mean to estimate  $V = \mathbb{E}[G]$ :

$$\hat{V}_n = \frac{1}{n} \sum_{k=1}^n G_k$$

- Do not need any information about how a dice works!
- Law of large numbers:  $\hat{V}_n \to \mathbb{E}[G]$  as  $n \to \infty$ .

# **Example: Throwing two dice**



Using the estimate

$$\hat{V}_n = \frac{1}{n} \sum_{k=1}^n G_k$$

| Trail | n = 1 | n = 10 | n = 100 | n = 1000 |
|-------|-------|--------|---------|----------|
| 1.    | 12.00 | 7.10   | 6.94    | 7.07     |
| 2.    | 9.00  | 7.20   | 7.30    | 6.92     |
| 3.    | 7.00  | 7.50   | 6.75    | 6.95     |
| 4.    | 4.00  | 4.90   | 7.02    | 7.01     |
| 5.    | 5.00  | 8.10   | 7.62    | 7.09     |

### Bias and variance



- Let  $\hat{\theta}_n$  be an estimate of  $\theta$  using n random samples.
- Since the samples are random,  $\hat{\theta}_n$  is a stochastic variable.
- ullet We can thus talk about the expected value and variance of  $\hat{\theta}_n$ .
- Bias: (Unbiased if bias= 0)

$$\mathsf{Bias}(\hat{\theta}_n) = \mathbb{E}[\hat{\theta}_n] - \theta.$$

Variance:

$$Var(\hat{\theta}_n) = \mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2]$$

• The mean squared error (MSE):

$$\mathsf{MSE}(\hat{\theta}_n) = \mathbb{E}[(\hat{\theta}_n - \theta)^2] = \mathsf{Var}(\hat{\theta}_n) + \mathsf{Bias}(\hat{\theta}_n)^2$$

• Consistent if  $\hat{\theta}_n \to \theta$  as  $n \to \infty$ .

# Example: Throwing two dice



- We want to estimate  $V = \mathbb{E}[G]$ .
- We use observations  $G_1, \ldots, G_n$  to form the estimate

$$\hat{V}_n = \frac{1}{n} \sum_{k=1}^n G_k.$$

• Bias:

$$\operatorname{Bias}(\hat{V}_n) = \mathbb{E}[\hat{V}_n] - V = \mathbb{E}\left[\frac{1}{n}\sum_{k=1}^n G_k\right] - V = \frac{1}{n}\sum_{k=1}^n \mathbb{E}[G] - V = \mathbb{E}[G] - V = 0$$

so the estimate is unbiased for all n.

Variance:

$$\operatorname{Var}(\hat{\theta}_n) = \mathbb{E}[(\hat{V}_n - \mathbb{E}[\hat{V}_n])^2] = \frac{35}{6n} \approx \frac{5.83}{n}$$

• So, as  $n \to \infty$  the variance goes to zero.

### Incremental updates



If we have already computed

$$\hat{V}_{n-1} = \frac{1}{n-1} \sum_{j=1}^{n-1} G_j,$$

and I get one more observation  $G_n$ :

$$\hat{V}_n = \frac{1}{n} \sum_{j=1}^n G_j = \frac{1}{n} \left( G_n + \sum_{j=1}^{n-1} G_j \right) \\
= \frac{1}{n} \left( G_n + (n-1)\hat{V}_{n-1} \right) = \hat{V}_{n-1} + \frac{1}{n} (G_n - \hat{V}_{n-1}).$$

ullet So, we can start from  $\hat{V}=0$ , n=0 and for each new observation G do

$$n \leftarrow n + 1$$
  
 $\hat{V} \leftarrow \hat{V} + \frac{1}{n}(G - \hat{V}).$ 

Fast update, and don't need to remember all old observations!

### **Incremental updates**



$$\underbrace{\hat{V}}_{\text{New Estimate}} \leftarrow \underbrace{\hat{V}}_{\text{Old Estimate}} + \underbrace{\alpha_n}_{\text{Step size}} \left[ \underbrace{G}_{\text{Target}} - \underbrace{\hat{V}}_{\text{OldEstimate}} \right]$$

- In each step, move the estimate a bit closer to the observed "target".
- For empirical mean the step size  $\alpha_n = \frac{1}{n}$ .
- ullet For i.i.d observations of G this converge to  $\mathbb{E}[G]$  if

$$\sum_{n=1}^{\infty} \alpha_n = \infty, \quad \sum_{n=1}^{\infty} \alpha_n^2 < \infty.$$

- Non-stationary problems: Can use a constant  $\alpha \in (0,1)$  to "forget" old observations. Variance do not go to zero, but can adjust to changing probabilities.
- Extreme cases:
  - $\alpha = 0$  gives  $\hat{V} \leftarrow \hat{V}$ . We learn nothing.
  - $\alpha=1$  gives  $\hat{V} \leftarrow G$ . We forget all past observations!

# Monte-Carlo Prediction

### **Monte-Carlo Prediction**



- Lets consider episodic tasks (that terminate after a finite number of steps).
- **Goal:** Learn  $v_{\pi}(s)$  from experience under  $\pi$ ,

$$S_0, A_0, R_1, S_1, A_1, R_2, \ldots, S_T \sim \pi.$$

The return:

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T.$$

• The value function:

$$v_{\pi}(s) = \mathbb{E}_{\pi}\left[G_t|S_t=s\right]$$

- Monte-Carlo: Estimate  $v_{\pi}(s)$  using the *empirical mean* return of many episodes instead of the *expected* return.
- With MC, we do not need to know what p(s', r|s, a) is!

### First-visit vs every-visit



#### First-visit

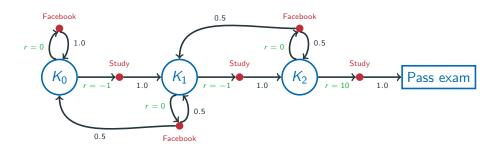
- 1. Sample an episode,  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$  using policy  $\pi$ .
- 2. **The first** time-step t that state s is visited, add  $G_t$  to Returns(s).
- 3. Let V(s) = average(Returns(s)).
- 4. Go back to step 1.

### **Every-visit**

- 1. Sample an episode,  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$  using policy  $\pi$ .
- 2. **Every** time-step t that state s is visited, add  $G_t$  to Returns(s).
- 3. Let V(s) = average(Returns(s)).
- 4. Go back to step 1.

# **Example: Study or Facebook?**





### **Trajectory:**

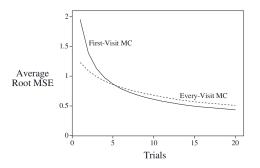
$$S_1 = K_0, R_2 = -1, S_2 = K_1, R_3 = -1, S_3 = K_2, R_4 = 0,$$
  
 $S_4 = K_1, R_5 = -1, S_5 = K_2, R_6 = 10, S_6 = Pass$ 

- First visit: For  $K_1$  only use  $G_2 = -1 + 0\gamma 1\gamma^2 + 10\gamma^3$ .
- Every visit: For  $K_1$  use  $G_2$  and  $G_4 = -1 + 10\gamma$ .

# Properties of MC



- + Consistent:  $V(s) \rightarrow v_{\pi}(s)$  as  $N(s) \rightarrow \infty$ .
- + First-visit MC is **unbiased**. (Every-visit can be biased).
- +/- Does not make use of the Markov-property.
  - Generally high variance, and reducing it may require a lot of experience.
  - Must wait until end of episode to compute  $G_t$  and update V.



"Reinforcement Learning with Replacing Eligibility Trace", Singh and Barto, 1996.

### Incremental updates



- In the MC method we compute the average of all observed returns  $G_t$  seen in each state.
- Incremental updates:
  - 1. Collect a trajectory  $S_0, R_1, S_1, R_2, \dots, S_T$  following policy  $\pi$ .
  - 2. For (the first-visit/every-visit)  $S_t$  compute  $G_t$  and let

$$N(S_t) \leftarrow N(S_t) + 1$$
  
 $V(S_t) \leftarrow V(S_t) + \alpha_n(G_t - V(S_t)).$ 

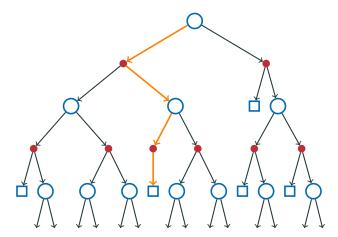
- For empirical mean:  $\alpha_n = \frac{1}{N(S_t)}$ .
- For e.g. non-stationary environment we may instead use a constant  $\alpha$ .

# Monte Carlo Backup Diagram



$$V(S_t) \leftarrow V(S_t) + \frac{1}{N(S_t)} (G_t - V(S_t))$$

The observations  $G_t$  can be computed from experience.



# Monte-Carlo vs Dynamic Programming



$$v_{\pi}(s) = \mathbb{E}_{\pi}[G_t|S_t = s] = \mathbb{E}_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1})|S_t = s]$$

### **Dynamic Programming:**

$$V(s) \leftarrow \mathbb{E}_{\pi}[R_{t+1} + \gamma V(S_{t+1}) | S_t = s].$$

- **Bootstrapping:** Each new estimate is based on a previous estimate.
- Computes expectation exactly, but estimate since based on estimate  $V(S_{t+1})$ .
- Needs model p(s', r|s, a) to compute expectation.

#### Monte Carlo:

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

- We do not bootstrap, since we use the full return  $G_t$ .
- Is an estimate because we use empirical mean of  $G_t$ , and not  $\mathbb{E}_{\pi}[G_t|S_t=s]$ .
- No model needed, since samples  $G_t$  can be computed from experience.

Can we combine bootstrapping and learning from experience?

Temporal-Difference (TD) Learning

# Temporal-Difference (TD) Learning



$$egin{aligned} v_{\pi}(s) &= \mathbb{E}_{\pi}[G_{t}|S_{t} = s] \ &= \mathbb{E}_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1})|S_{t} = s] \end{aligned}$$

• MC: We use the target  $G_t$ :

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

• **TD:** We use the *TD-target*  $R_{t+1} + \gamma V(S_{t+1})$ :

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

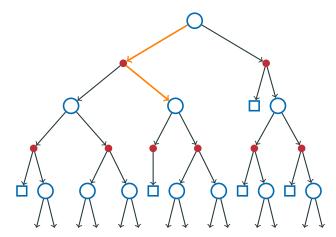
- Often called TD(0), since it is a special case of TD( $\lambda$ ) (with  $\lambda = 0$ ).
- TD **bootstraps**: The new estimate  $V(S_t)$  is based on the estimate  $V(S_{t+1})$ .

### **TD Backup Diagram**



$$V(S_t) \leftarrow V(S_t) + \alpha \left( R_{t+1} + \gamma V(S_{t+1}) - V(S_t) \right)$$

We do not have to complete a full episode to make the update!



### **TD-learning**



- Initialize the estimate V (e.g. V(s) = 0 for all s)
- Start in some state *S*.
- Loop
  - 1. Take action A according to policy  $\pi(a|S)$ .
  - 2. Observe reward R and new state S'.
  - 3.  $V(S) \leftarrow V(S) + \alpha [R + \gamma V(S') V(S)]$
  - 4.  $S \leftarrow S'$ .
- (If the task is episodic, we would have to re-run the above loop for several episodes).

**Note:** We do not have to complete the episode before we start learning, and we can even learn in continuing tasks.



$$egin{aligned} v_{\pi}(s) &= \mathbb{E}_{\pi}[G_t|S_t = s] \ &= \mathbb{E}_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1})|S_t = s] \end{aligned}$$

### The MC-target $G_t$

- **Unbiased** estimate of  $v_{\pi}(S_t)$ .
- Not based on previous estimates. (No bootstrapping)

# The "true TD-target": $R_{t+1} + \gamma v_{\pi}(S_{t+1})$

- **Unbiased** estimate of  $v_{\pi}(S_t)$ .
- Cannot be computed when we do not know  $v_{\pi}(S_{t+1})$ .

# The TD-target $R_{t+1} + \gamma V(S_{t+1})$

- **Biased** estimate of  $v_{\pi}(S_t)$ .
- Based on old estimate of  $V(S_{t+1})$  (Bootstrapping)

#### TD vs MC



### Monte Carlo (MC):

- Only in episodic environments.
- High variance, zero bias.
- Converges to  $v_{\pi}(s)$  (if  $\alpha$  decrease with a suitable rate)
- (good convergence properties even for function approximation)
- Not very sensitive to initial conditions.
- Usually more efficient in non-MDP environments.

### Temporal Differences (TD):

- Both episodic and continuing environments.
- Low variance, but some bias.
- Converges to  $v_{\pi}(s)$  (if  $\alpha$  decrease with a suitable rate)
- (but do not always converge with function approximation)
- More sensitive to initial conditions.
- Usually more efficient in MDP environments.

#### Batch MC and TD



- MC and TD converge as experience  $\to \infty$ .
- Lets assume that we only have a finite number of episodes of experience:

$$S_0^1, A_0^1, R_1^1, \cdots, S_T^1$$
  
 $\vdots$   
 $S_0^K, A_0^K, R_1^K, \cdots, S_T^K$ 

- Batch learning: Repeatedly apply MC or TD to the episodes we have.
- ullet With constant (but small) lpha both methods converge, but to different estimates.
- Batch MC: Minimize the mean-squared error  $\sum (G_t - V(S_t))^2$  on training data.
- Batch TD: Equivalent to finding the maximum likelihood estimate of  $p(s', r|s, \pi(s))$ , and then computing the value function according to the estimated model.

# **AB Example**



Two states (A and B),  $\gamma = 1$ .

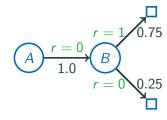
### **Experience:**

- B, 1
- B, 0

### Batch MC: Converges to

$$V(A) = 0, V(B) = 0.75$$

Batch TD: Converges to



and thus

$$V(A) = 0.75, V(B) = 0.75.$$

# Extensions\*

# *n*-step return and $TD(\lambda)$



• In MC we use the full return as target:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + R_T.$$

In TD, we only take one step and then use our previous estimate.
 That is the TD-target

$$G_{t:t+1} = R_{t+1} + \gamma V(S_{t+1}).$$

In the same way we could define a 2-step return

$$G_{t:t+2} = R_{t+1} + \gamma R_{t+2} + \gamma^2 V(S_{t+2})$$

and so on.

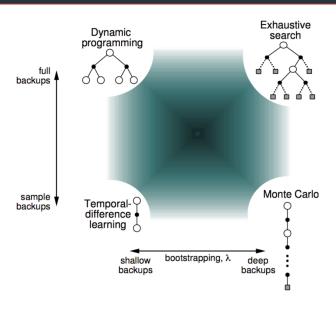
• In  $TD(\lambda)$  we average between different *n*-step returns and use

$$G_t^{\lambda} = (1 - \lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} G_{t:t+n} + \lambda^{T-t-1} G_t.$$

Note that  $\lambda=0$  gives the 1-step return, and  $\lambda=1$  gives MC.

### Unified view of RL





### **Summary**



- How to estimate  $v_{\pi}$  from experience using MC and TD(0).
- Both converge to  $v_{\pi}$  in tabular case (if  $\alpha$  decrease with a suitable rate).
- MC unbiased (first-visit), high variance.
- TD biased, relatively low variance.
- TD usually more efficient in MDP-environments.

#### Next:

- How to find a good policy when p(s', r|s, a) is unknown.
- You can now also look at Tinkering Notebook 3a.